

# **11868 LLM Systems**

# **Deep Learning Framework**

# **Design**

Lei Li



**Carnegie Mellon University**  
Language Technologies Institute

# Recap

- Learning algorithm for Neural Network
  - stochastic gradient descent
- Computation Graph
  - topological traversal along the DAG
- Auto Differentiation
  - building backward computation graph

# Today's Topic

- ➡ • How to design a deep learning framework
  - Design ideas in TensorFlow
    - Abadi et al., “TensorFlow: A System for Large-Scale Machine Learning”, OSDI 2016

# Need for DL Programming System

- Deep learning already claiming big successes
- Huge need for high-productivity tools for developing machine learning solutions for various applications
- Instead of writing cuda and differentiation code for each specific model

# Deep Learning Programming Framework

- Open source library for machine learning computation using data flow graphs
- TensorFlow is an interface for expressing machine learning algorithms, and an implementation for executing such algorithms
- PyTorch is a programming framework for tensor computation, deep learning, and auto differentiation


# TensorFlow

- Key idea: express a numeric computation as a computation graph
  - following what we described in last lecture
- Graph nodes are **operations** with any number of inputs and outputs
- Graph edges are **tensors** which flow between nodes
  - tensor: multidimensional array

# Programming Model

Computation graph in tensorflow

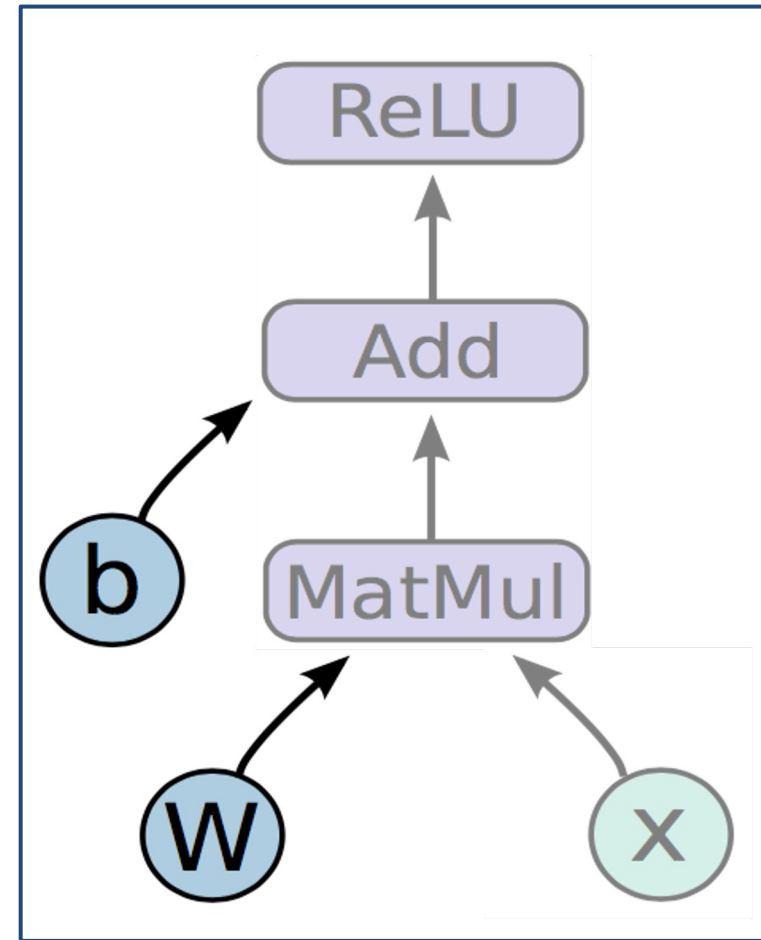
$$h = \text{RELU}(Wx + b)$$



# Variables

$$h = \text{ReLU}(Wx + b)$$

- **Variables** are stateful nodes which output their current value.
- State is retained across multiple executions of a graph
- mostly parameters

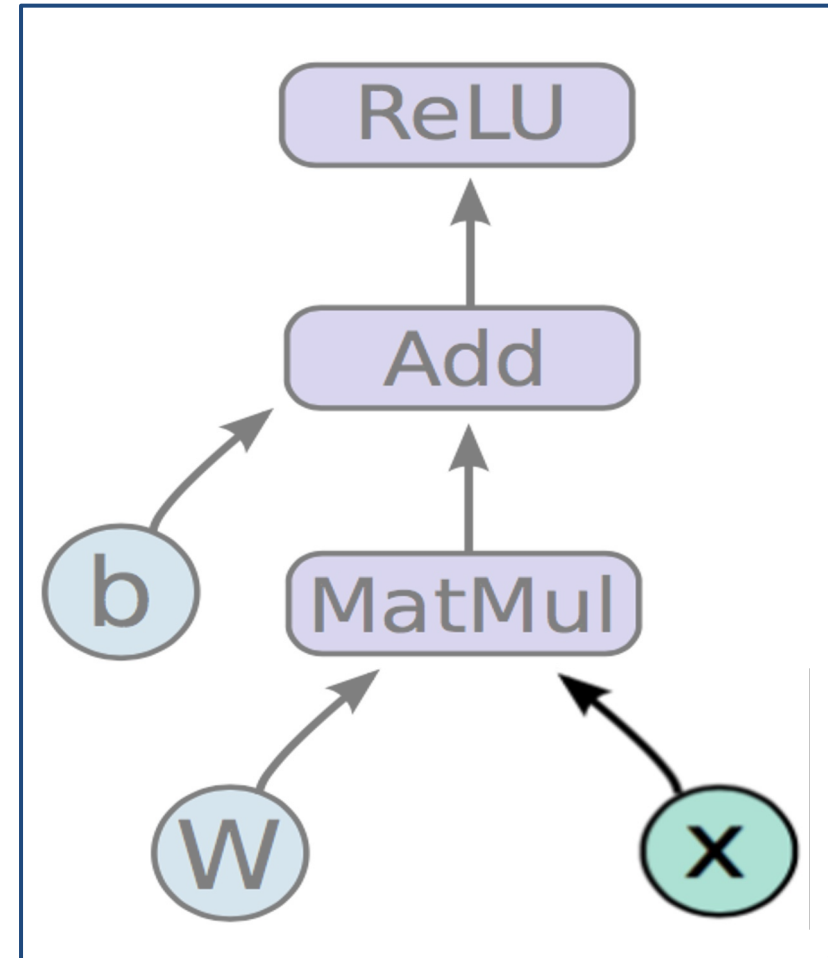




# Placeholders

$$h = \text{ReLU}(Wx + b)$$

- **Placeholders** are nodes whose value is fed in at execution time
- Inputs, Labels, ...



# Mathematical Operations


$$h = \text{ReLU}(Wx + b)$$

**MatMul**: Multiply two matrices

**Add**: Add elementwise

**ReLU**: Activate with elementwise rectified linear function

$$\text{ReLU}(x) = \begin{cases} 0, & x \leq 0 \\ x, & x > 0 \end{cases}$$



# Programming the Graph

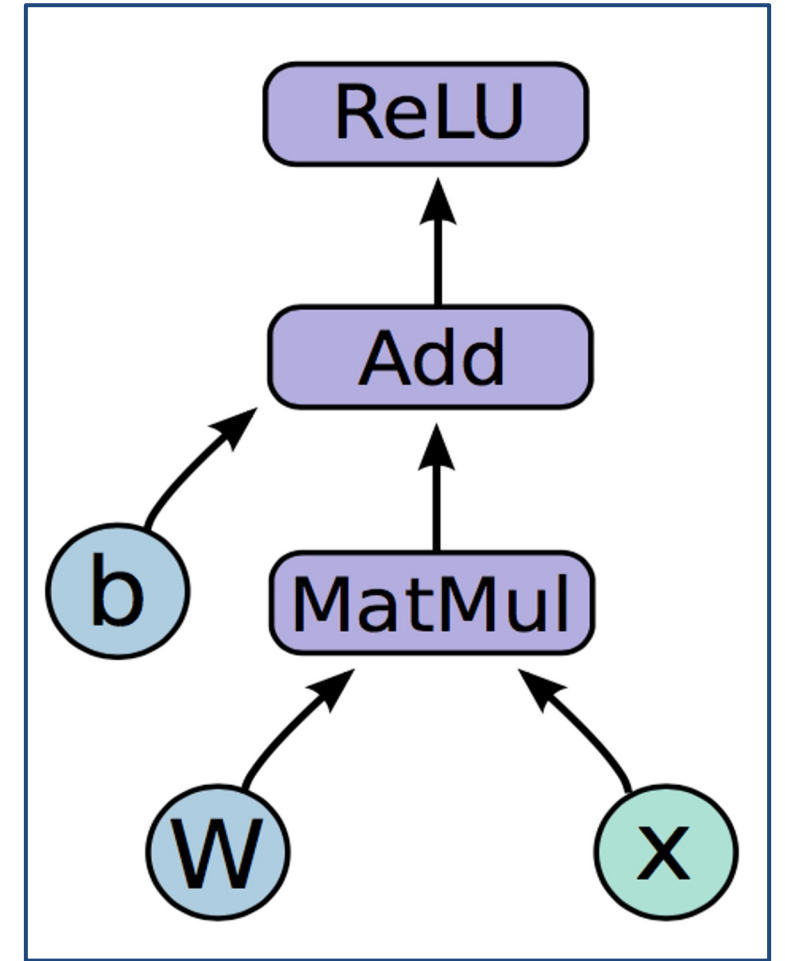
```
import tensorflow as tf

b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))

x = tf.placeholder(tf.float32, (1, 784))

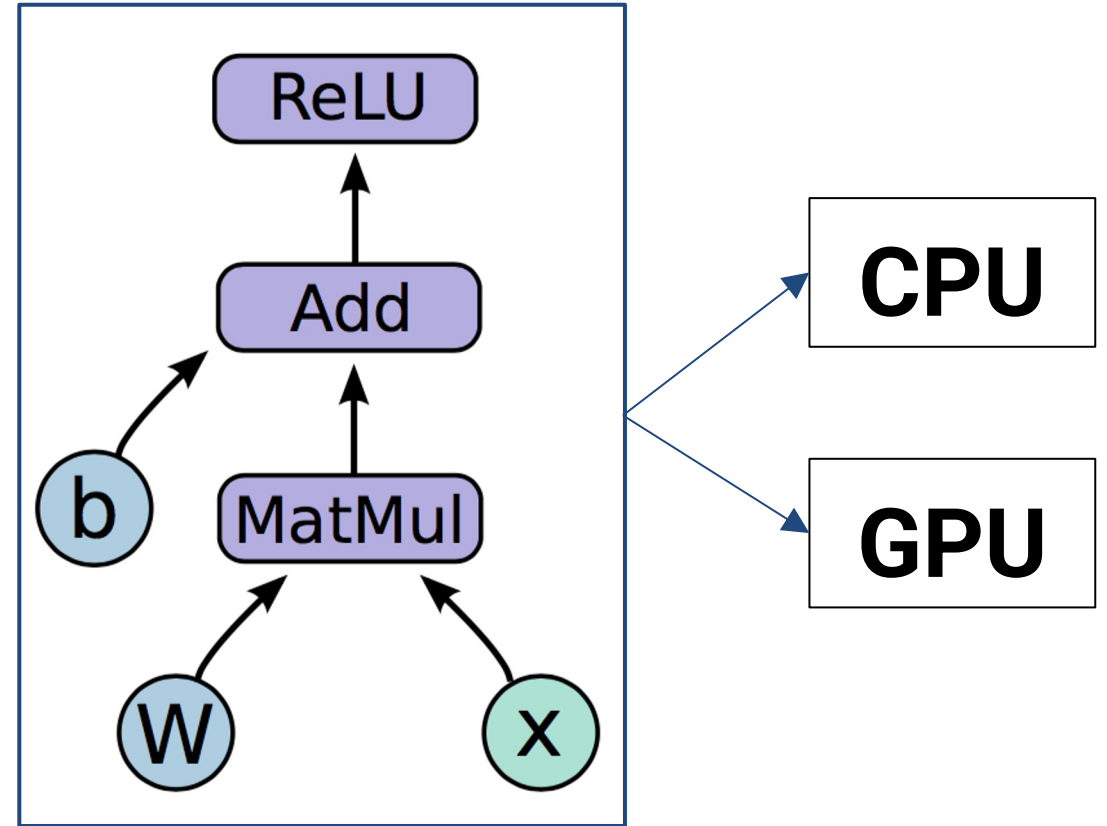
h = tf.nn.relu(tf.matmul(x, W) + b)
```

$$h = \text{RELU}(Wx + b)$$



# Running the Graph

Deploy graph with a **session**: a binding to a particular execution context (e.g. CPU, GPU)





```
1 import tensorflow as tf
2
3 with tf.Session() as sess:
4     # Phase 1: constructing the graph
5     a = tf.constant(15, name="a")
6     b = tf.constant(5, name="b")
7     prod = tf.multiply(a, b, name="Multiply")
8     sum = tf.add(a, b, name="Add")
9     res = tf.divide(prod, sum, name="Divide")
10
11     # Phase 2: running the session
12     out = sess.run(res)
13     print(out)
```

# Defining Loss

- Use placeholder for labels
- Build loss node using labels and prediction

```
prediction = tf.nn.softmax(...) #Output of neural network
```

```
label = tf.placeholder(tf.float32, [100, 10])
```

```
cross_entropy = -tf.reduce_sum(label *  
tf.log(prediction), axis=1)
```

# Gradient Computation

```
train_step =  
tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```

- `tf.train.GradientDescentOptimizer` is an `Optimizer` object
- `tf.train.GradientDescentOptimizer(lr).minimize(cross_entropy)` adds optimization operation to computation graph
- TensorFlow graph nodes have attached gradient operations
- Gradient with respect to parameters computed with Auto Differentiation (recall previous lecture)

# Core TensorFlow Constructs

- All nodes return tensors, or higher-dimensional matrices
- How a node computes is indistinguishable to TensorFlow
- You are metaprogramming - constructing the graph for the real computation. No computation occurs yet!



# Implementing Graph Nodes

# Design Principles

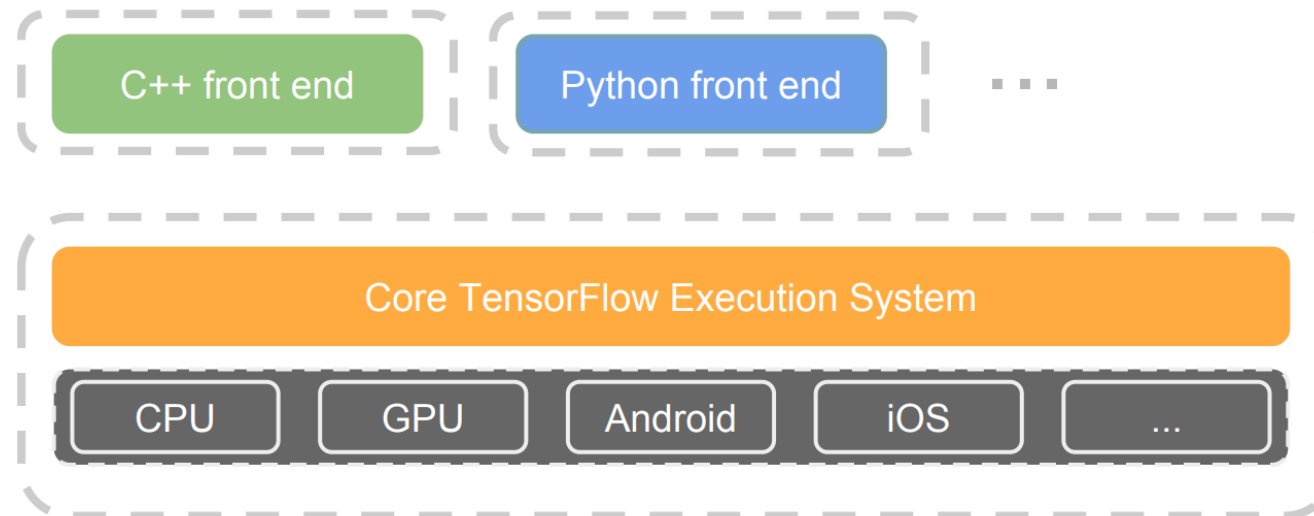
- Dataflow graphs of primitive operators
- Deferred execution (two phases)
  1. Define program i.e., symbolic dataflow graph w/ placeholders
  2. Executes optimized version of program on set of available devices

# Dynamic Flow Control

- Problem: support ML algos that contain conditional and iterative control flow, e.g.
  - Recurrent Neural Networks (RNNs) and LSTMs
  - Autoregressive decoder
- Solution: Add conditional (if statement) and iterative (while loop) programming constructs

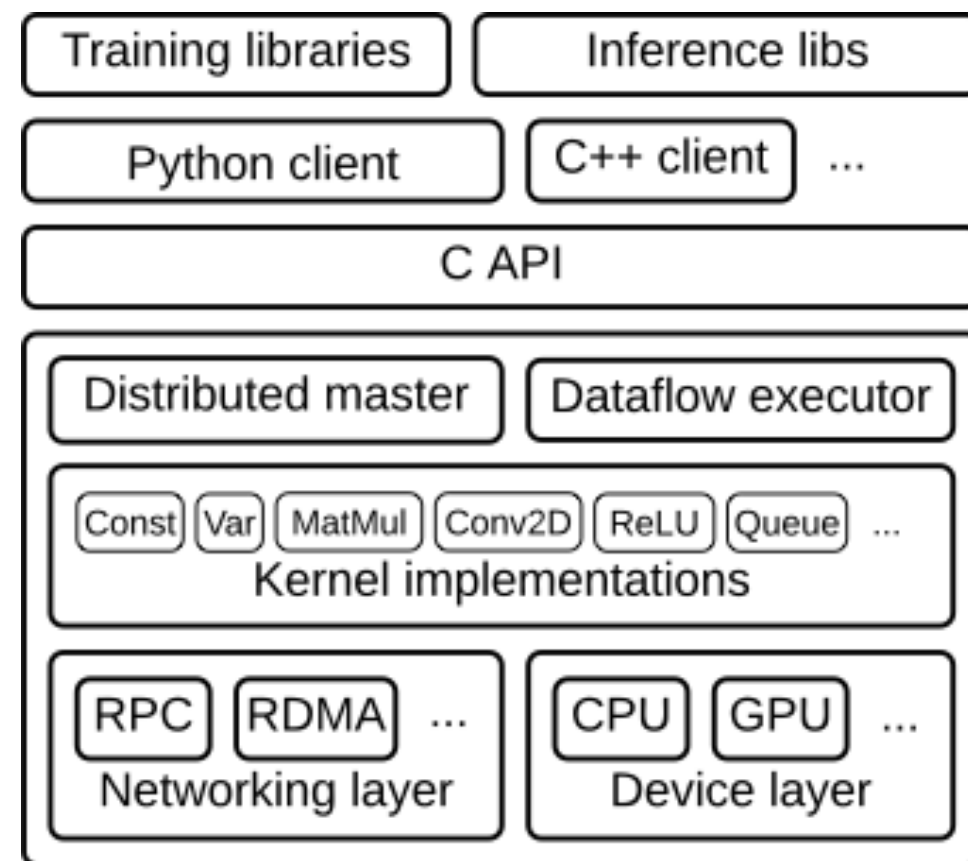
# TensorFlow Architecture

- Core in C++
  - Very low overhead
- Different front ends for specifying/driving the computation
  - Python and C++, easy to add more




# TensorFlow Implementation

- Semi-interpreted
- Call to kernel per primitive operation
- Can batch operations with custom C++
- Basic type-safety within dataflow graph (error at graph construction time)




# Key Components

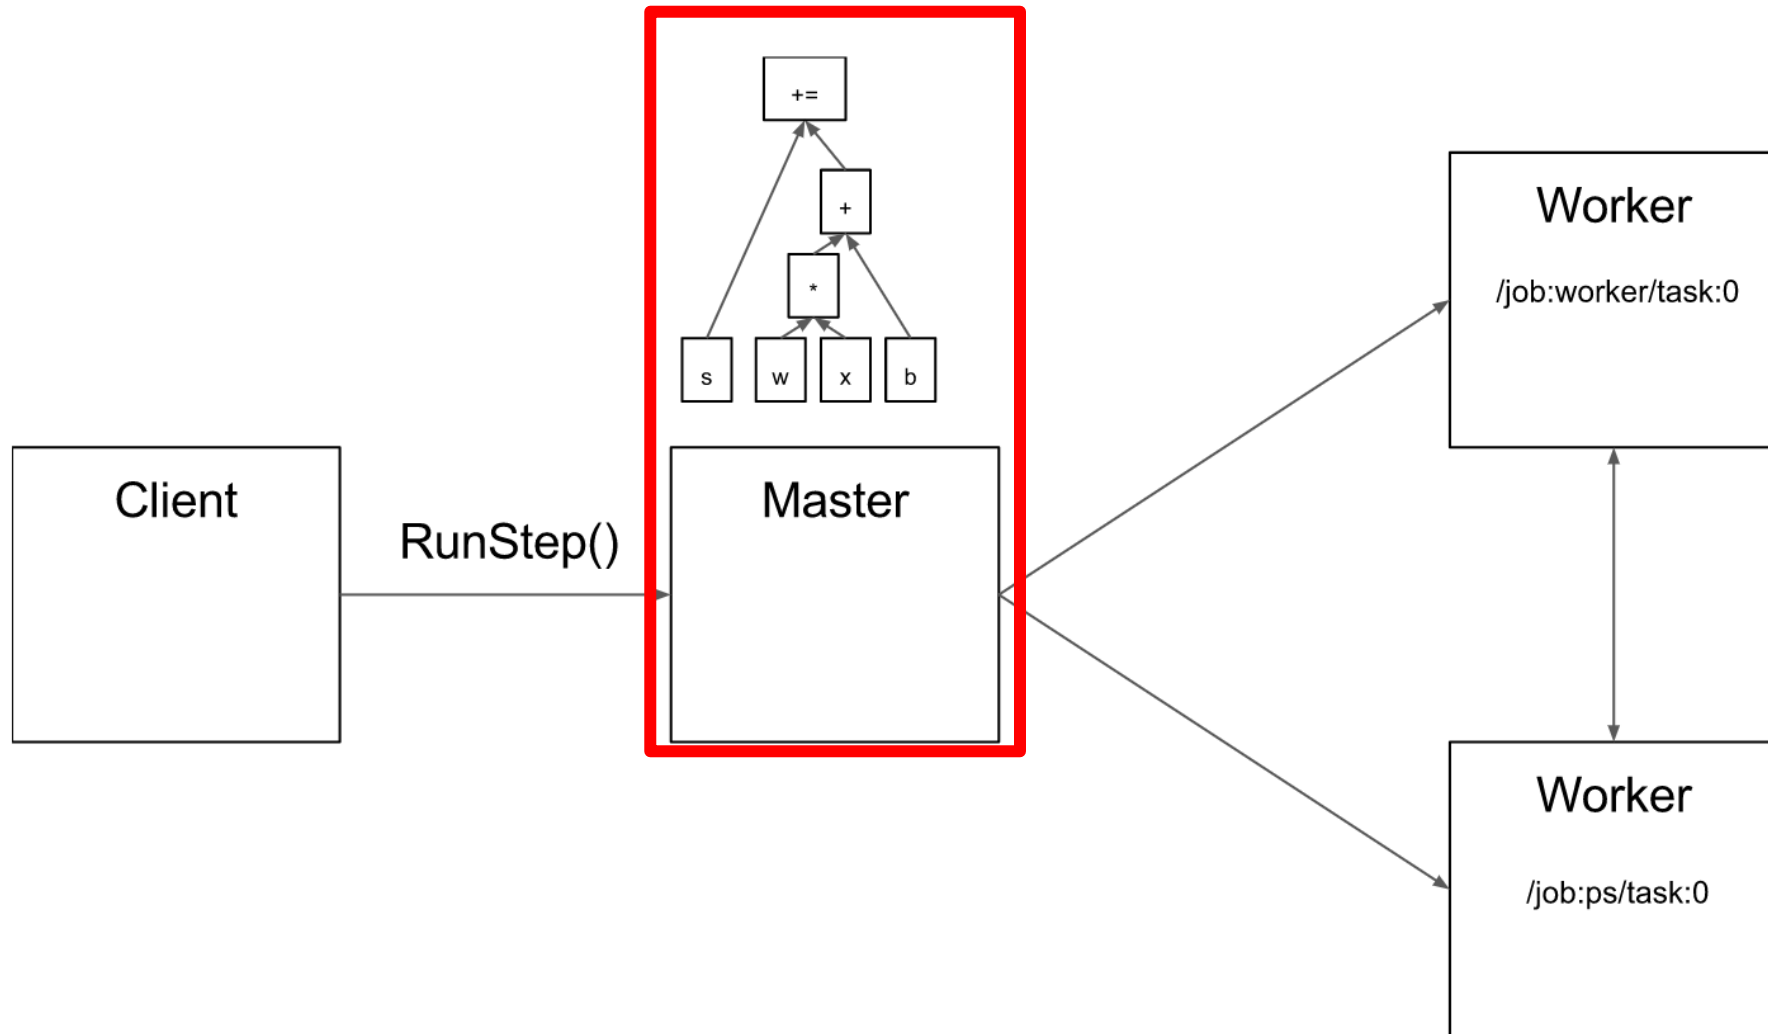
- Similar to MapReduce, Apache Hadoop, Apache Spark, ...



# Client



# Master

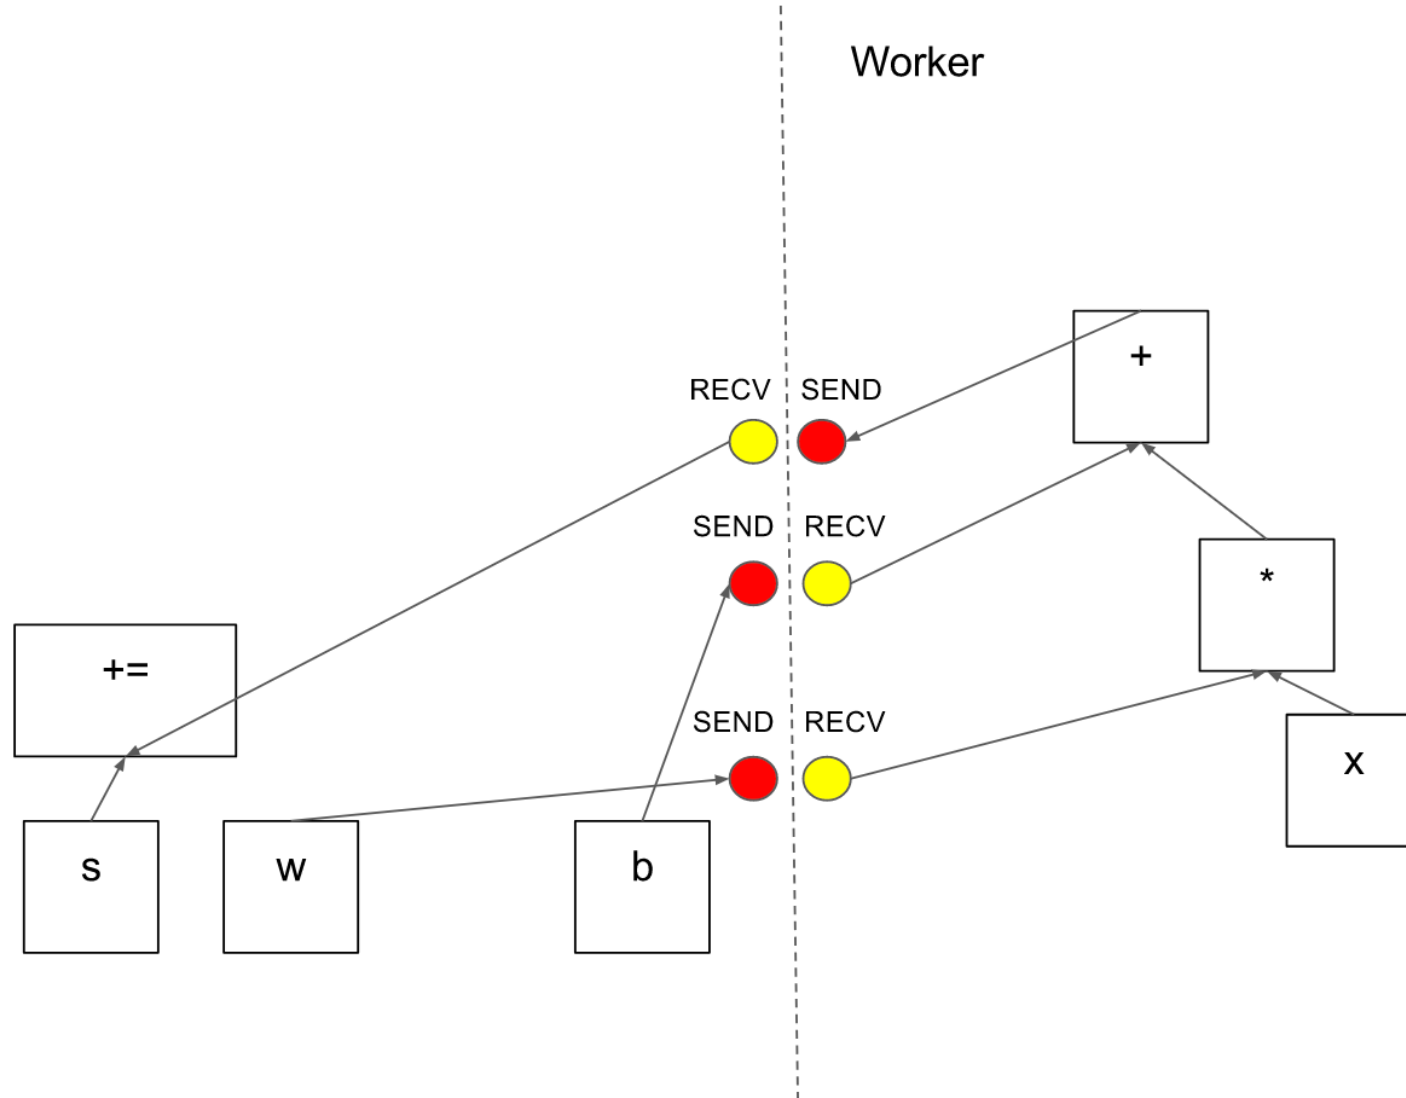





# Computation Graph Partition

PS

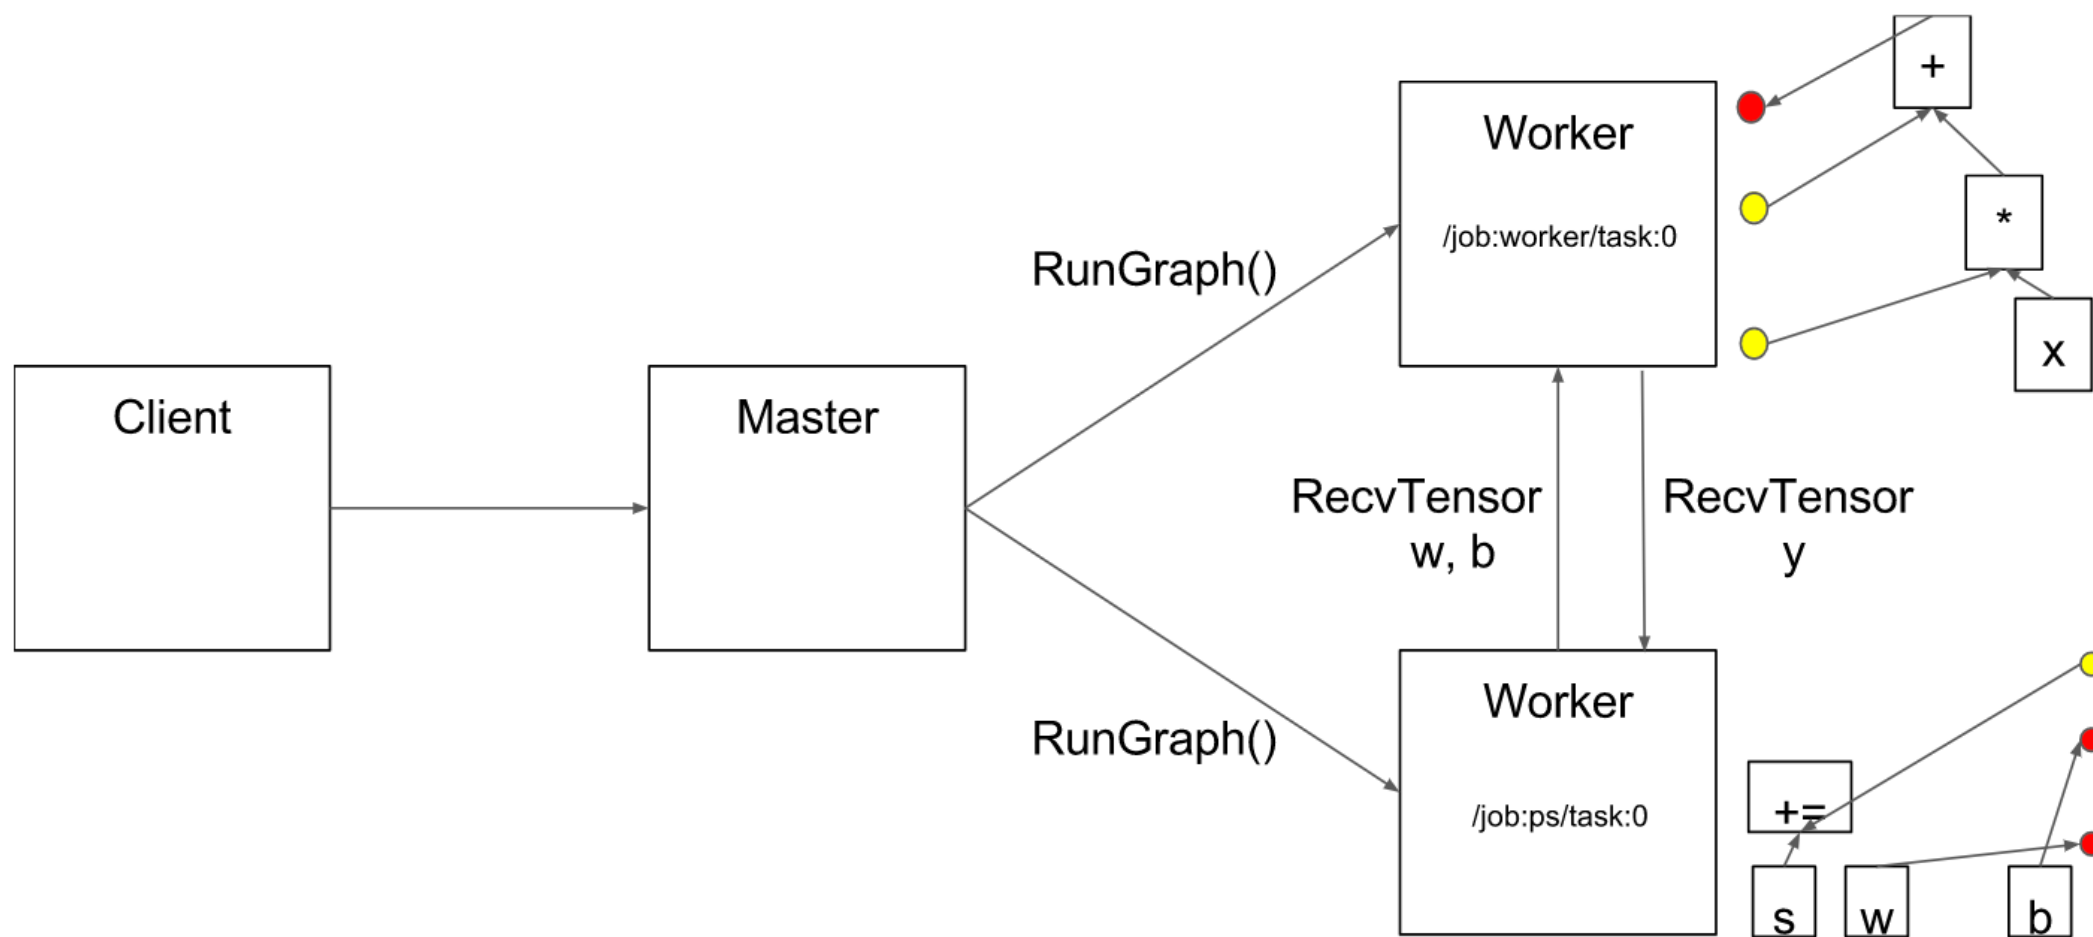
Worker



# Computation Graph Partition



# Execution



# Synchronous vs Asynchronous

- Determined by node: Queue nodes used for barriers
- Synchronous nearly as fast as asynchronous
- Default model is asynchronous

# Fault Tolerance

- Assumptions:
  - Fine grain operations: “It is unlikely that tasks will fail so often that individual operations need fault tolerance” ;-)
  - “Many learning algorithms do not require strong consistency”
- Solution: user-level checkpointing (provides 2 ops)
  - `save()`: writes one or more tensors to a checkpoint file
  - `restore()`: reads one or more tensors from a checkpoint file



# Performance

- Single Node

Library	Training step time (ms)			
	AlexNet	Overfeat	OxfordNet	GoogleNet
Caffe [38]	324	823	1068	1935
Neon [58]	87	<b>211</b>	<b>320</b>	<b>270</b>
Torch [17]	<b>81</b>	268	529	470
TensorFlow	<b>81</b>	279	540	445

# Performance

- Distributed Throughput



# Summary

- Key Contributions
  - Programmability
  - Accessibility / ease of use
  - Richness of Libraries
  - Ready-made community



# MiniTorch Code Explanation

# Reading for Next Class

- Attention is all you need. 2017